




kubernetes

# Kubernetes training Itenium

Jan Wielemans  
[jan.wielemans@devgem.be](mailto:jan.wielemans@devgem.be)



kubernetes



docker

Evening 1  
24/11/2022



kubernetes



kubernetes

Evening 2  
22/12/2022



Debugging



Istio



kubernetes

# Specific questions day 1

- **Docker logs**

[Local File logging driver | Docker Documentation](#)

<https://tecadmin.net/truncate-docker-container-logfile/>

- **Docker image labeling schema**

[Label Schema | Bringing Structure To Labels \(label-schema.org\)](#)



# Recap



- **Docker**
  - CLI `docker`
  - Pull and create images
  - Run and configure containers
  - `docker-compose`
- **Kubernetes**
  - Lens
  - CLI `kubectl`
  - Pod



kubernetes



# Service

## Service


- abstraction
- loadbalancing
- type
- based on selectors

## Pod

my-app

```
apiVersion: v1
kind: Service
metadata:
  name: my-app-service
spec:
  selector:
    app: my-app-pod
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

# Service types







kubernetes

# Demo time

## Service



# Ingress

\*.goca-training.be  
aaa.bbb.ccc.ddd

NGINX

Loadbalancer

Ingress Service

Ingress Controller

myapp.goca-training.be

Service

Deployment

my-app

app1.goca-training.be

Service

Deployment

app2

app2.nogietsanders.com

Service

Deployment



app3

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-app-ingress
spec:
  rules:
    - host: myapp.goca-training.be
      http:
        paths:
          - pathType: Prefix
            path: /
            backend:
              service:
                name: my-app-service
                port:
                  number: 80
  ingressClassName: nginx
```



kubernetes





# Namespace

```
apiVersion: v1
kind: Namespace
metadata:
  name: demo
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  namespace: demo
  name: demo-backend
  labels:
    app: demo-backend
spec:
  ...
```



kubernetes

## What?

Logical division of the cluster in virtual sub-clusters (that can communicate internally and can have RBAC)

## Why?



Division by application or environment (e.g. test, staging) or whatever you want

## How?

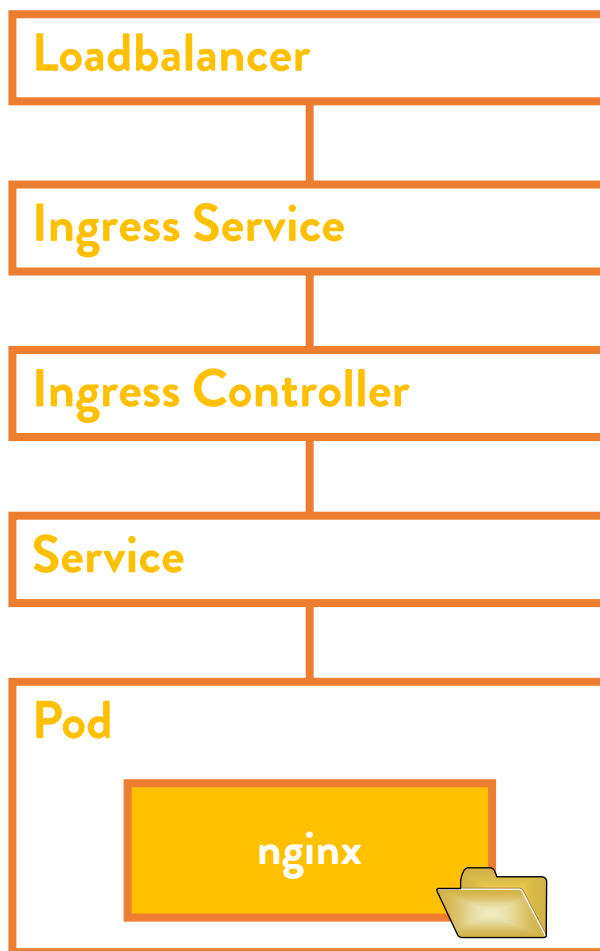
Assign each resource to a namespace (except nodes and persistent volumes, default is default, don't use kube-system, kube-public or kube-node-lease)



## Demo time

- Ingress
- Multi api resource
- Namespaces
- TLS



# Volumes





# Volumes



kubernetes

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: nginx-pvc
spec:
  storageClassName: do-block-storage
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
```



```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: nginx-pod
spec:
  containers:
    - name: nginx-container
      image: nginx:1.23.1
      volumeMounts:
        - mountPath: /usr/share/nginx/html
          name: nginx-volume
  volumes:
    - name: nginx-volume
      persistentVolumeClaim:
        claimName: nginx-pvc
```



kubernetes

# Demo time

## Volumes



# Deployment

## Deployment

- metadata (e.g. labels)
- specs (replicas, selectors)
- pod template
- scaling
- rolling updates

demo-db-mongodb



```
apiVersion: apps/v1
kind: Deployment
metadata:
  namespace: demo
  name: demo-db-mongodb
  labels:
    app: demo-db-mongodb
spec:
  replicas: 1
  selector:
    matchLabels:
      app: demo-db-mongodb
  template:
    metadata:
      namespace: demo
      labels:
        app: demo-db-mongodb
    spec:
      containers:
        - name: demo-db-mongodb
          image: mongo:latest
```



kubernetes

## Demo time


- Deployment
- Scale & loadbalance
- Rolling updates
- Probe
- Bringing it all together
- Intro debugging




# Horizontal Pod Autoscaler



kubernetes







# Cluster Autoscaler



kubernetes

The screenshot shows the GitHub interface for the `kubernetes/autoscaler` repository. The browser address bar shows the URL `https://github.com/kubernetes/autoscaler/blob/master/cluster-autoscaler/README.md`. The repository page includes navigation links (Features, Business, Explore, Marketplace, Pricing), a search bar, and options to sign in or sign up. The repository statistics show 61 watches, 592 stars, and 276 forks. The file view shows the `cluster-autoscaler/README.md` file, which is 126 lines long and 6.86 KB in size. The README content includes the title **Cluster Autoscaler**, an **Introduction** section, and a list of conditions under which the autoscaler adjusts the cluster size.


Cluster Autoscaler

## Introduction

Cluster Autoscaler is a tool that automatically adjusts the size of the Kubernetes cluster when:

- there are pods that failed to run in the cluster due to insufficient resources.
- some nodes in the cluster are so underutilized, for an extended period of time, that they can be deleted and their pods will be easily placed on some other, existing nodes.

## FAQ/Documentation



**docker**




**kubernetes**




**kubernetes**

CLI's:


- **docker**
- **kubectl**




**Pod**




**Ingress**




**Volumes**




**Deployment**




**Secret**




**Namespace**



**Horizontal Pod Autoscaler**



**Service**

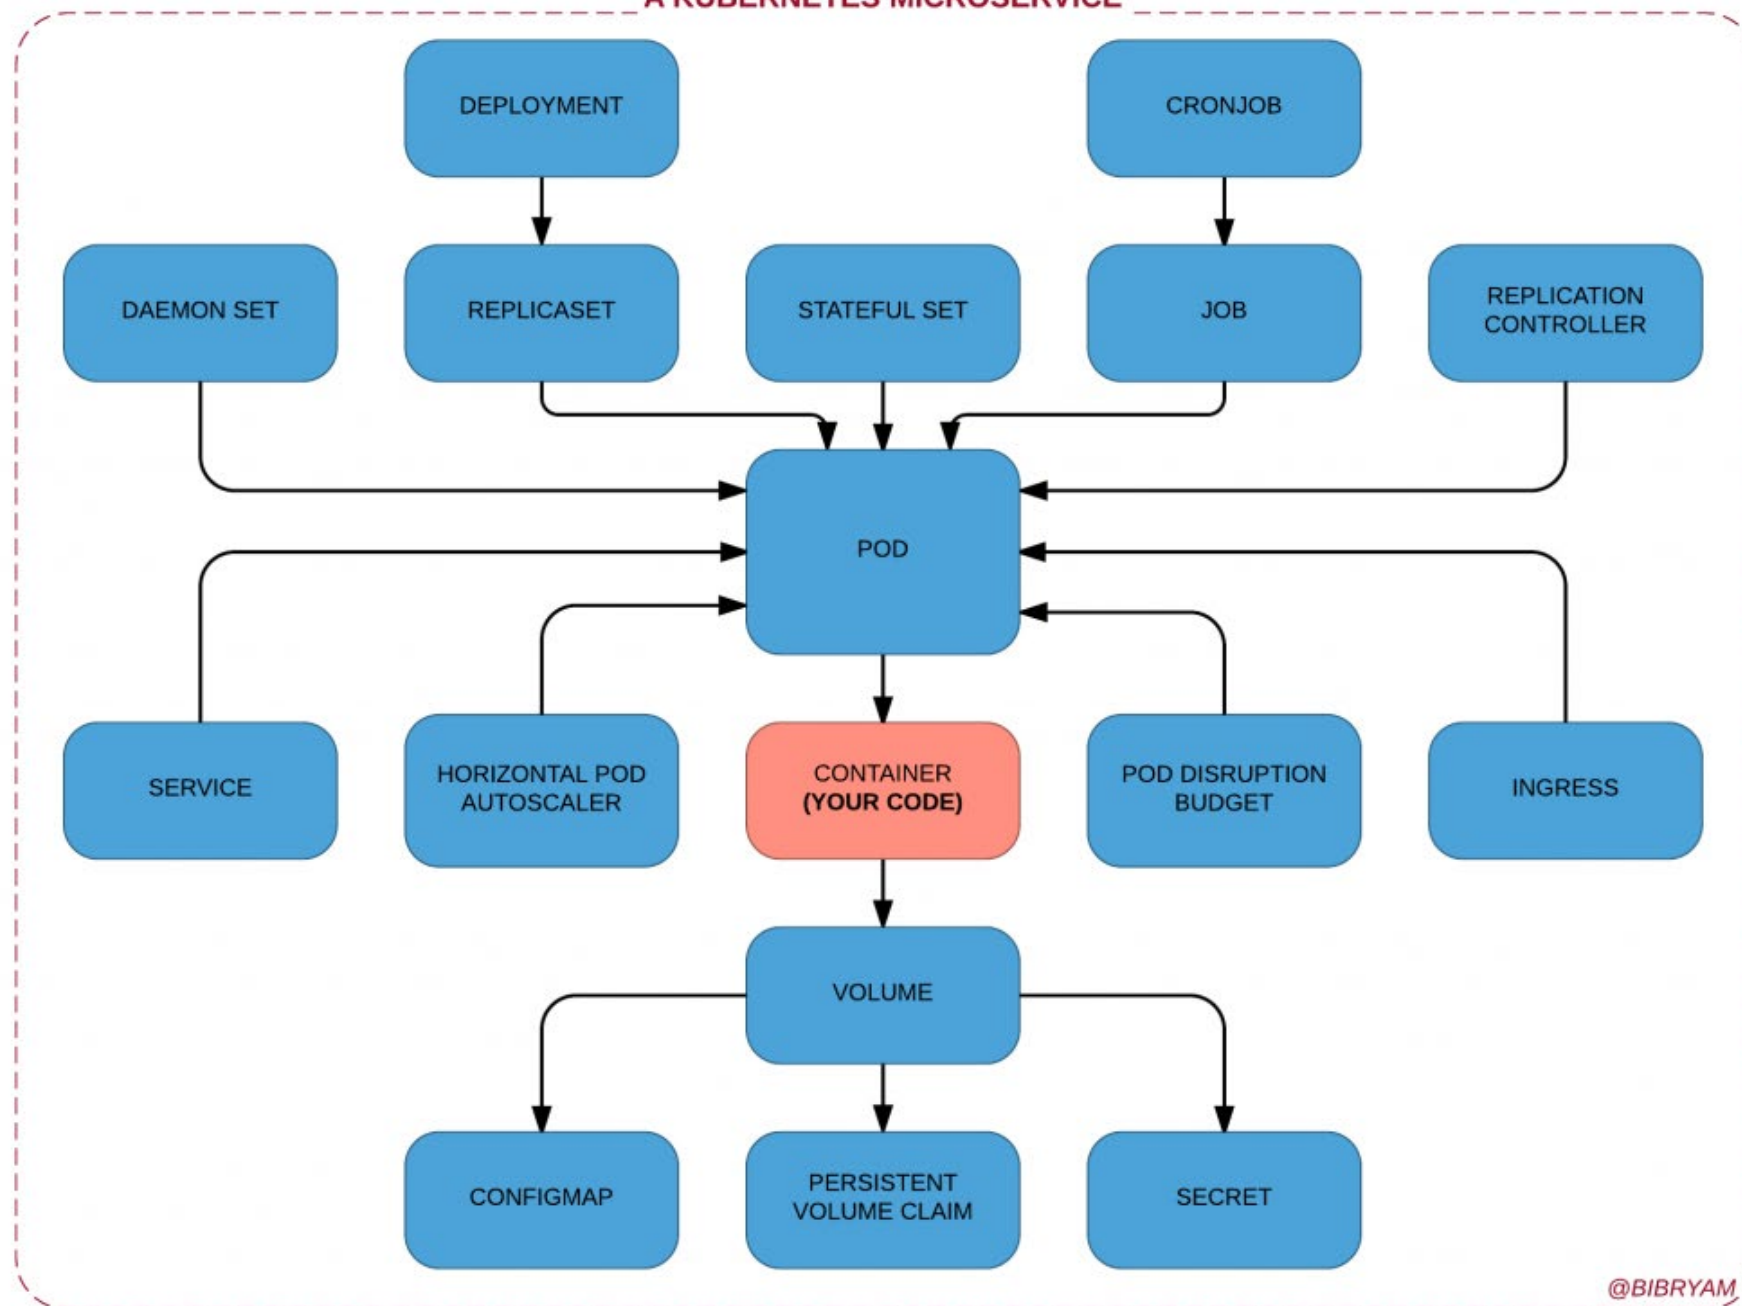


**Cluster Autoscaler**

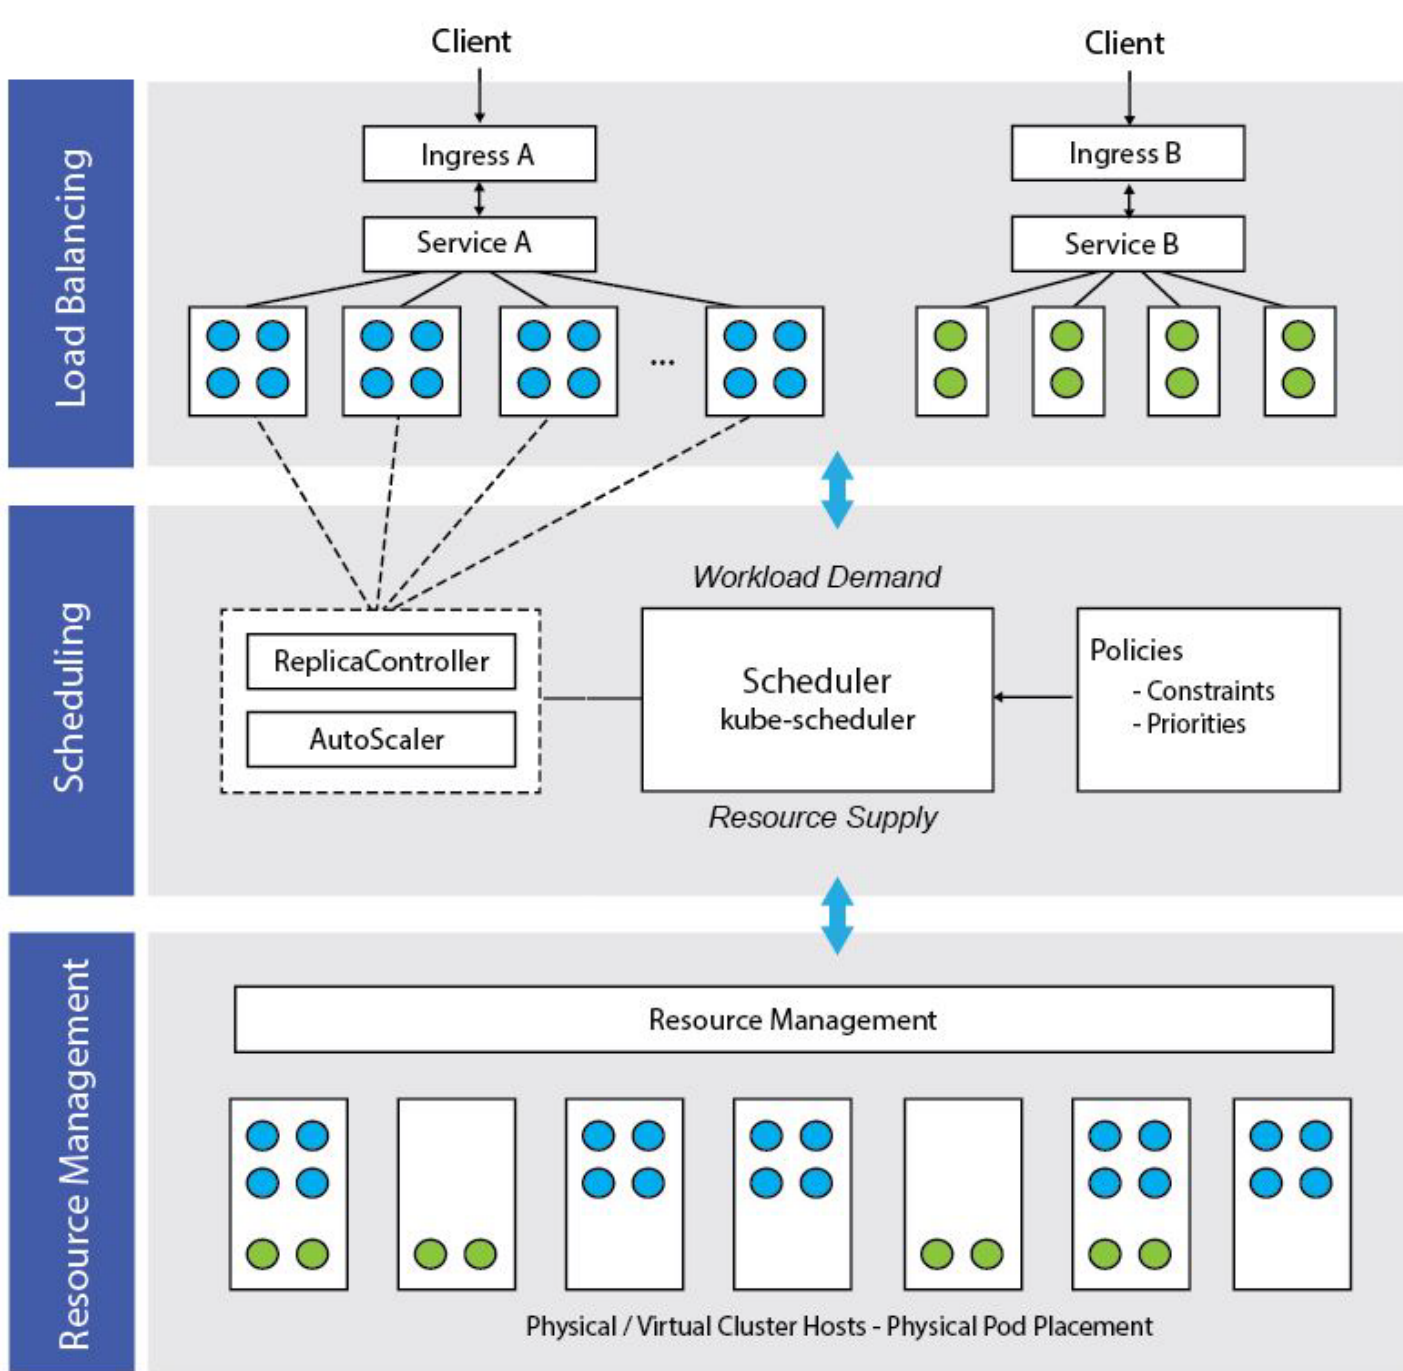


kubernetes

## A KUBERNETES MICROSERVICE



@BIBRYAM





kubernetes



# Helm is a package manager for Kubernetes



## Isn't the Docker image the package?

Yes, but typically when deploying an entire application it can consist of several containers (the container landscape of your application) with specific properties and “wiring”.

## Aren't the Kubernetes .yaml files then the way to package the application?

Yes, but typically you would want to have different sets of files with different configurations for different environments (e.g. for test and production environments you would have different namespaces, different url's, different db connectionstrings, different scaling, etc.).

## So how does Helm help?

Helm is

- a CLI
- a convention for how to organize .yaml files (a folder structure that can be zipped called a “chart”)
- a templating language to make your .yaml files configurable

# Helm illustration



```
wordpress/
Chart.yaml      # A YAML file containing information about the chart
LICENSE         # OPTIONAL: A plain text file containing the license for the chart
README.md      # OPTIONAL: A human-readable README file
values.yaml     # The default configuration values for this chart
values.schema.json # OPTIONAL: A JSON Schema for imposing a structure on the values.yaml file
charts/        # A directory containing any charts upon which this chart depends.
crds/          # Custom Resource Definitions
templates/     # A directory of templates that, when combined with values,
               # will generate valid Kubernetes manifest files.
templates/NOTES.txt # OPTIONAL: A plain text file containing short usage notes
```

```
demo:
  environment: staging
  url: demo-api-staging.goca-training.be
```

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  namespace: { .Values.demo.environment }
  name: demo-backend
spec:
  rules:
    - host: { .Values.demo.url }
      http:
        paths:
          - pathType: Prefix
            path: /
            backend:
              service:
                name: demo-backend
                port:
                  number: 80
            ingressClassName: nginx
```

```
helm install charts/demo --set demo.environment=production --set demo.url=demo-api.goca-training.be
```



kubernetes



Applications

Kubernetes

Developers

Company

# Helm Charts

Find your favorite application in our catalog and launch it.

Learn more about the [benefits of the Bitnami Application Catalog](#).





kubernetes

# Demo time

- Deploying a *MongoDB* cluster in your Kubernetes cluster with Helm



kubernetes



# Istio

# Istio is a Service Mesh

## Why

A service mesh enables developers to separate and manage service-to-service communications in a dedicated infrastructure layer. As the number of microservices involved with an application increases, so do the benefits of using a service mesh to manage and monitor them.



## Concepts



### Traffic Management

Deploy capabilities like inter-service routing, failure recovery and load balancing.



### Observability

Provide an end-to-end view of traffic flow and service performance.



### Security

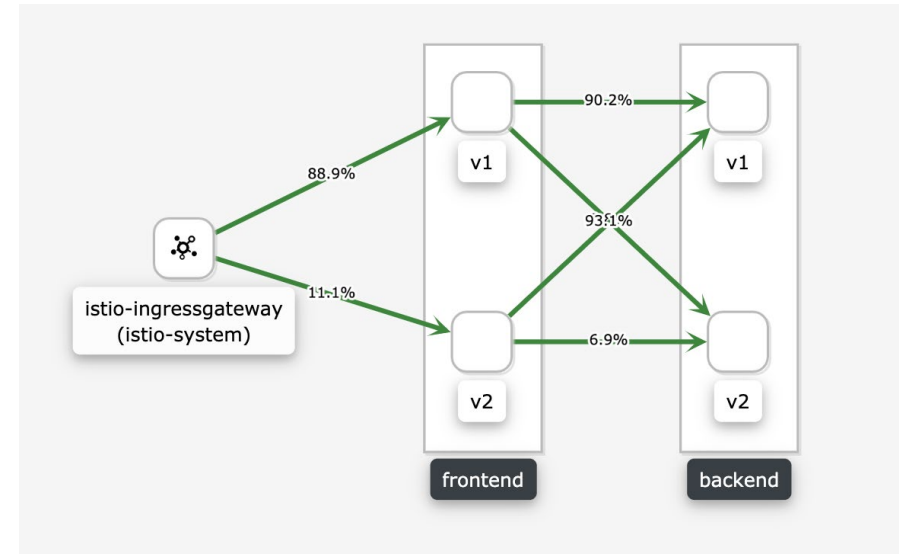
Engage encryption, role-based access, and authentication across services.

# A service mesh tackles microservice challenges

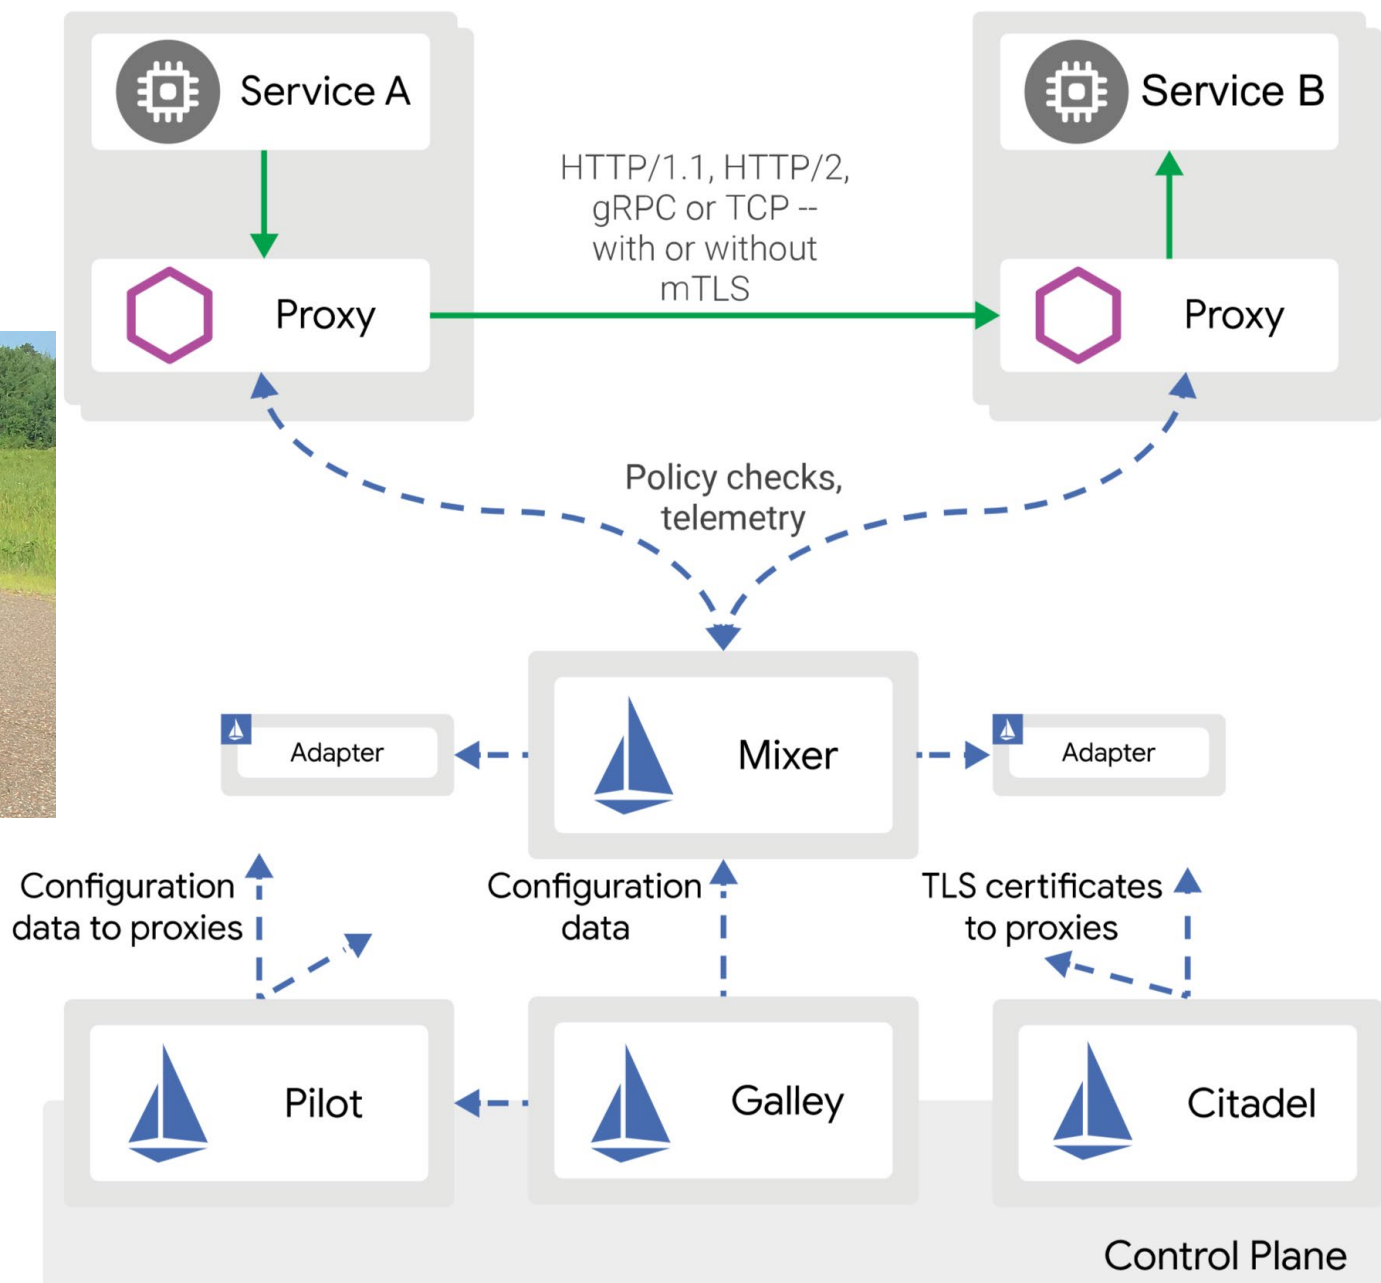


kubernetes

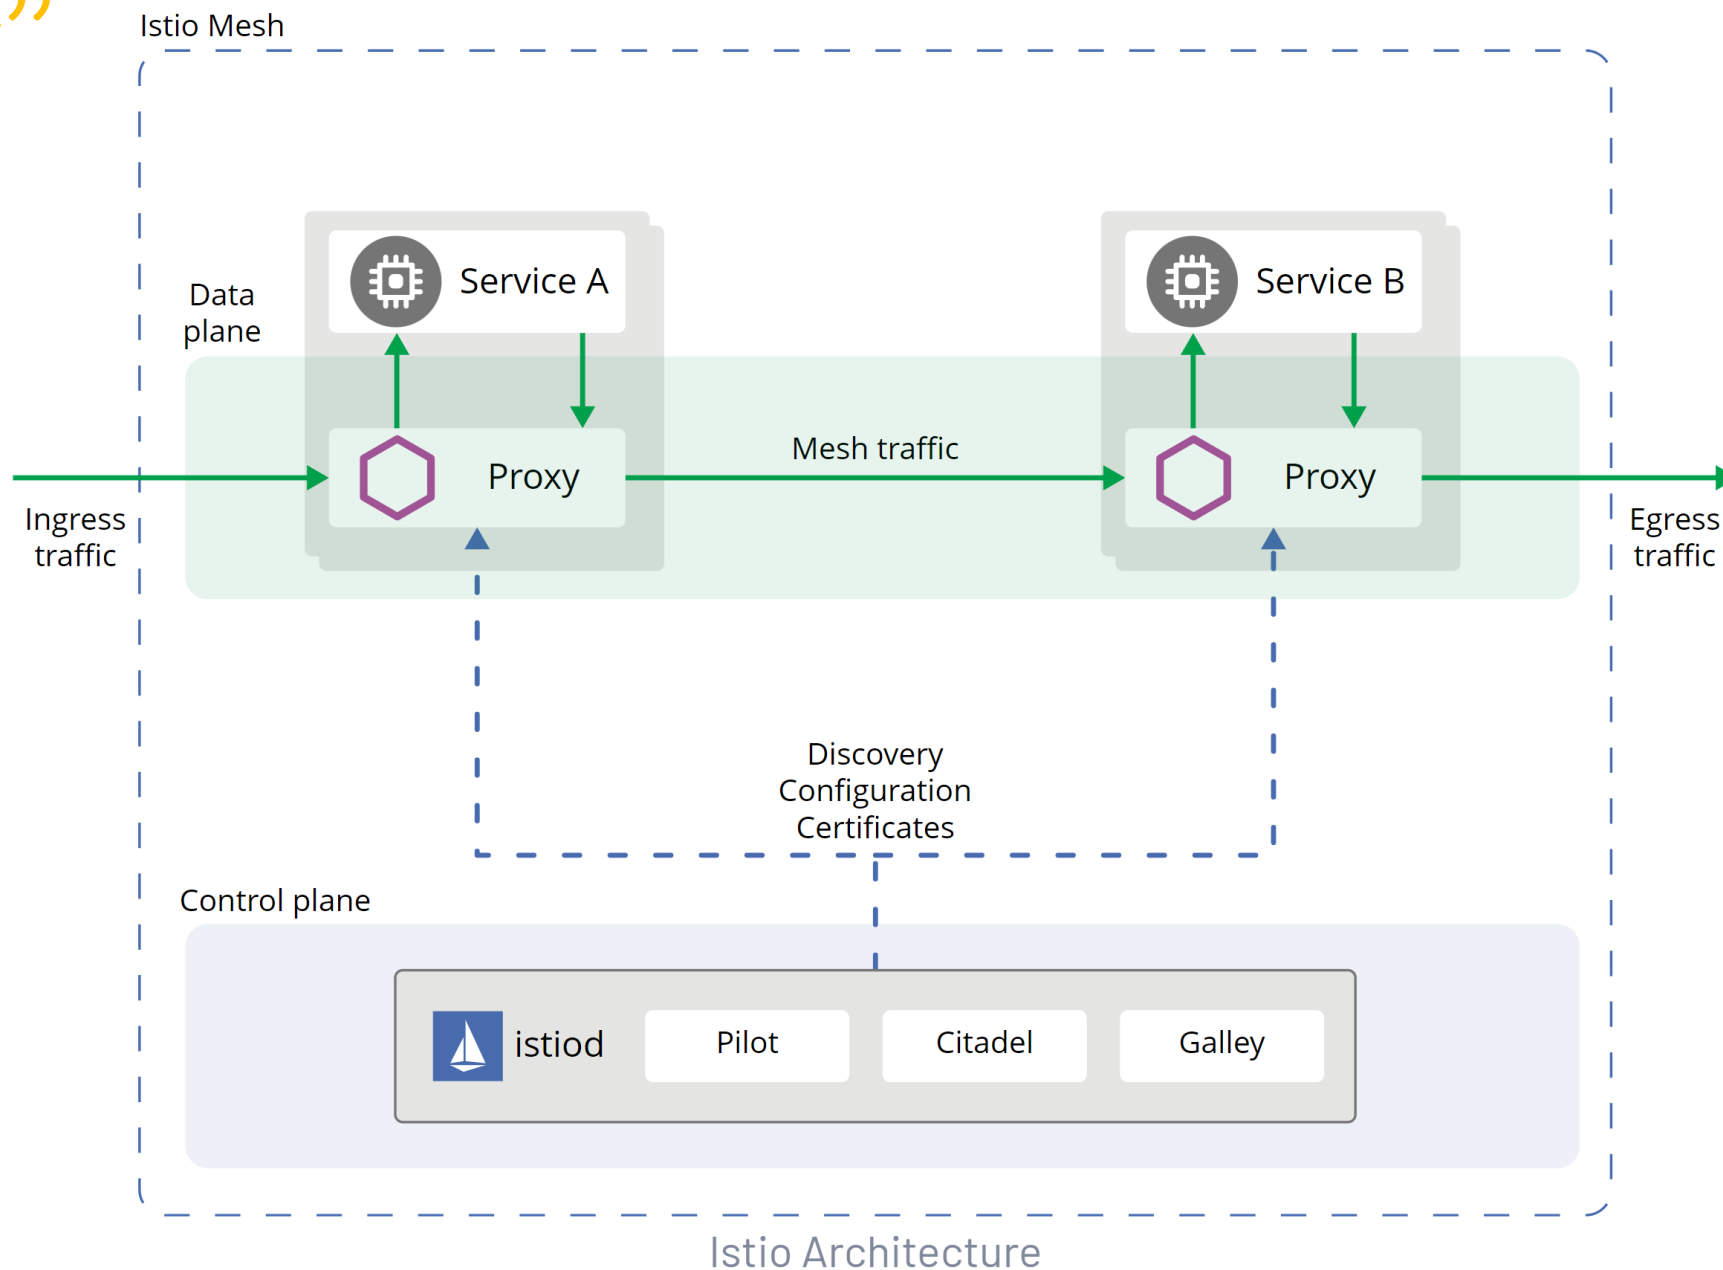
- Security
- Canary deployments
- A/B testing
- Retries
- Rate limiting
- Fault injection
- Policy management
- Telemetry



# “Classic” Istio



# “New” Istio





kubernetes

# Scratch the surface...

- Ingress
- Traffic shaping
- Chaos testing
- Observability



kubernetes

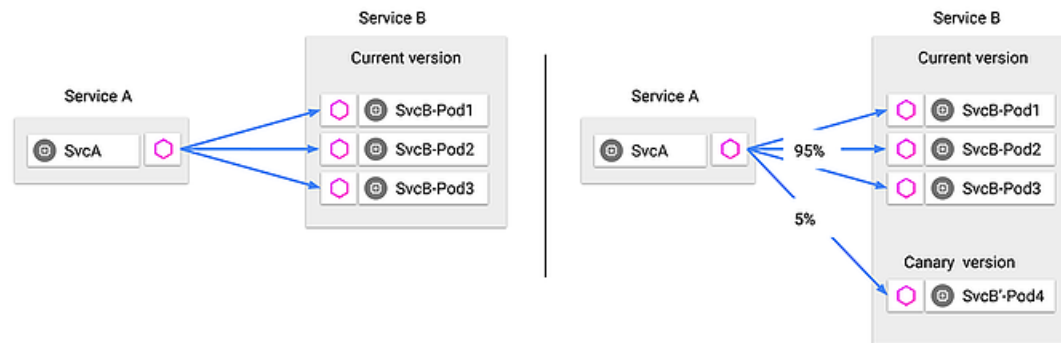
# Demo time

- Istio installation
- Istio ingress

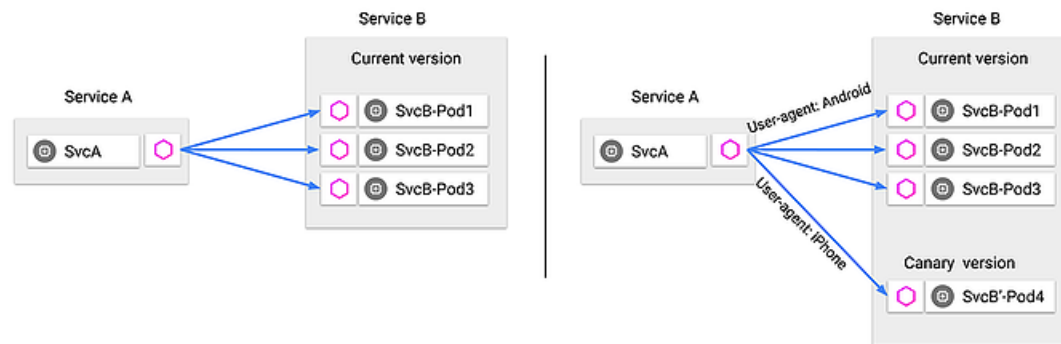


# Traffic shaping

## Canary release



**Traffic splitting decoupled from infrastructure scaling** - proportion of traffic routed to a version is independent of number of instances supporting the version



**Content-based traffic steering** - The content of a request can be used to determine the destination of a request

## Dark release

# Demo time

- Dark deployment
- Canary deployment

# Chaos testing



## Benefits of Chaos testing

Below are the key benefits of Chaos testing

Benefits of Chaos Testing	
Five-Nines availability	One of the key benefits of chaos engineering is the very high availability of the system for its end users. Five-Nine availability means the system is up 99.999%. This means there are very less chances of system outages.
Financial profits	Even a very small outage can cause companies to lose millions of dollars. With chaos testing promising to keep the system up, companies are eying at increasing revenues.
Better disaster recovery plan in place	Chaos testing is a way to proactively eliminate, or at least reduce, the frequency and severity of any system disaster. The teams are more equipped to handle those, and therefore have better plans in place. The plans get better with more disasters avoided or recovered
Efficient coding	Since engineers know that their code will be tested for Chaos testing, they are challenged to write better codes to ensure the final system is as resilient as possible. They start thinking out of the box and bring innovative ideas into place.



kubernetes

# Demo time

- Fault injection

# Observability



## Demo time

- Install the observability stack
- Inspect the dashboards



kubernetes



# Istio

There's more...

[Istio / Tasks](#)



kubernetes



MicroK8s



**K3S**



<https://microk8s.io/docs/install-raspberry-pi>

<https://docs.k3s.io/advanced#raspberry-pi>






# Local Kubernetes DevOps pipeline

[Tilt | Kubernetes for Prod, Tilt for Dev](#)



kubernetes



docker

Evening 1  
24/11/2022



kubernetes



kubernetes

Evening 2  
22/12/2022



Debugging



Istio